

CS2109S: Introduction to AI and Machine Learning

Lecture 2: Solving Problems by Searching

25 August 2023

Admin: Assessments

- **Midterm (35%)**

- Date/Time: Friday, 6 October (Week 7), 10:00 AM
- Venue: Multipurpose Sports Hall (MPSH) (to be confirmed)

- **Final (35%)**

- Date/Time: Saturday, 18 November, 8:00 PM - Sunday, 19 November, 11:59 PM

Admin: Next Week Lecture

Next week (Friday, 1st September) is Polling Day

- There will be **no in-class lecture**
- Lecture will be recorded and published online
- You are expected to watch the lecture before the weekend is over

Admin: Tutorial

Tutorial starts next week!

- Tutorial allocations is available in Coursemology
- Tutorial swaps is allowed (see latest announcement)

Do not appeal through EduRec!

Recap

Credit: Cooking with AMC



- What is AI?
- History lesson
- **PEAS**: Performance Measure, Environment, Actuators, Sensors
- Properties of the task environment
 - Fully observable, deterministic, episodic, static, discrete, single-agent
- Agents
 - Common: reflex, model-based, goal-based, utility-based, learning

Outline

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Outline

- **Problem-solving agents**
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Problem-Solving Agents

When the correct action to take is not immediately obvious, an agent may need to **plan-ahead**: to consider a *sequence* of actions that form a path to a goal state. Such an agent is called a **problem-solving agent**, and the computational process it undertakes is called **search**.



Credit: Passport & Plates



Want to go to Bucharest, how?

Problem-Solving Process

- **Goal formulation**
- **Problem formulation**
 - Create an *abstract* model of the relevant parts of the world
- **Search**
 - Simulates sequence of actions using the model, search until goal is reached
- **Execution**
 - Execute actions in the solution, one at a time

We assume that the task environment is **fully observable** and **deterministic**.


Thus, the **solution** is a **fixed sequence of actions**.



Problem Formulation

Formally, a search problem can be formulated as follows.


- **States (state space).**
- **Initial state:** the initial state of the agent
- **Goal state(s)/test**
- **Actions:** things that the agent can do
- **Transition model:** what each action does
- **Action cost function:** the cost of performing an action



A sequence of action form a **path/trajectory**. **Solution** is a path to a goal.


Problem Formulation: Romania

- **States:** {at Sibiu, at Fagaras, ..., at Bucharest}
- **Initial state:** at Sibiu
- **Goal state(s)/test:** at Bucharest
- **Actions:** go to neighboring city x
- **Transition model:** move to target city
- **Action cost function:** distance




Problem Formulation: Vacuum World

- **States:** see image
- **Initial state:** any state
- **Goal state(s)/test:** all clean
- **Actions:** suck, left, right
- **Transition model:** see image
- **Action cost function:** 1/action



Problem Formulation: Grid World

- **States:** positions of agent
- **Initial state:** (0,1)
- **Goal state(s)/test:** (1,2)
- **Actions:** up, down, left, right
- **Transition model:** move if there is no wall, die on lava
- **Action cost function:** -1



Problem Formulation: 8-Puzzle

- **States:** location of tiles
- **Initial state:** see image
- **Goal state(s)/test:** see image
- **Actions:** move blank left, right, up, down
- **Transition model:** move blank
- **Action cost function:** 1/move

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8

Goal State

Problem Formulation: Robot Assembly

- **States:** robot joints positions and angles
- **Initial state:** item not assembled
- **Goal state(s)/test:** item assembled
- **Actions:** see image
- **Transition model:** move/rotate joint
- **Action cost function:** rotation/displacement



Outline

- Problem-solving agents
- **Search algorithms**
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Search Algorithms

A search algorithm takes in a search problem as input and returns a solution / failure. It is defined by the **order of node expansion**.

Evaluation criteria:

- **Time complexity:** number of nodes expanded
- **Space complexity:** maximum number of nodes in memory
- **Completeness:** does it return a solution if it exists?
- **Optimality:** does it always find the least-cost solution?

Measure: branching factor (b), depth (d), maximum depth (m)

Tree Search


Define order of node expansion

```
create frontier
insert initial state
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Representation Invariant

An **abstraction function** AF maps a real-world (concrete) state $c \in C$ to an **abstract representation** $a \in A$.



Representation invariant I is a function such that $I(a) = true$ for all legitimate representations $a \in A$. If $I(a) = true$, then $AF^{-1}(a) \in C$.

We need to make sure that our transition function satisfies the representation invariant i.e., produce valid states

Outline

- Problem-solving agents
- Search algorithms
- **Uninformed search algorithms**
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Uninformed Search Algorithms

An uninformed search algorithm is given no clue about how close a state is to the goal(s).

- Breadth-first search (BFS)
- Uniform-cost search
- Depth-first search (DFS)

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	$O((V + E)\log V)$
On Tree	BFS / DFS	$O(V)$
On DAG	Topological Sort	$O(V + E)$

CS2040S focus was on finding the shortest path (SSSP, APSP)

CS2109S: might not care about shortest path, or even the path


Breadth-first Search

```
create frontier : queue

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```



Breadth-first Search


```
create frontier : queue

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```



Queue:



Breadth-first Search

```
create frontier : queue  
  
insert initial state  
while frontier is not empty:  
    state = frontier.pop()  
    for action in actions(state):  
        next state = transition(state, action)  
  
        if next state is goal: return solution  
        frontier.add(next state)  
return failure
```

Queue:



Breadth-first Search

```
create frontier : queue




insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)

return failure
```

Queue:



Breadth-first Search

```
create frontier : queue
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```



```
    for action in actions(state):
```

```
        next state = transition(state, action)
```

```
        if next state is goal: return solution
```

```
        frontier.add(next state)
```

```
return failure
```



Queue:



Breadth-first Search

create **frontier** : **queue**

insert **initial state**

while **frontier** is not empty:

state = **frontier**.pop()


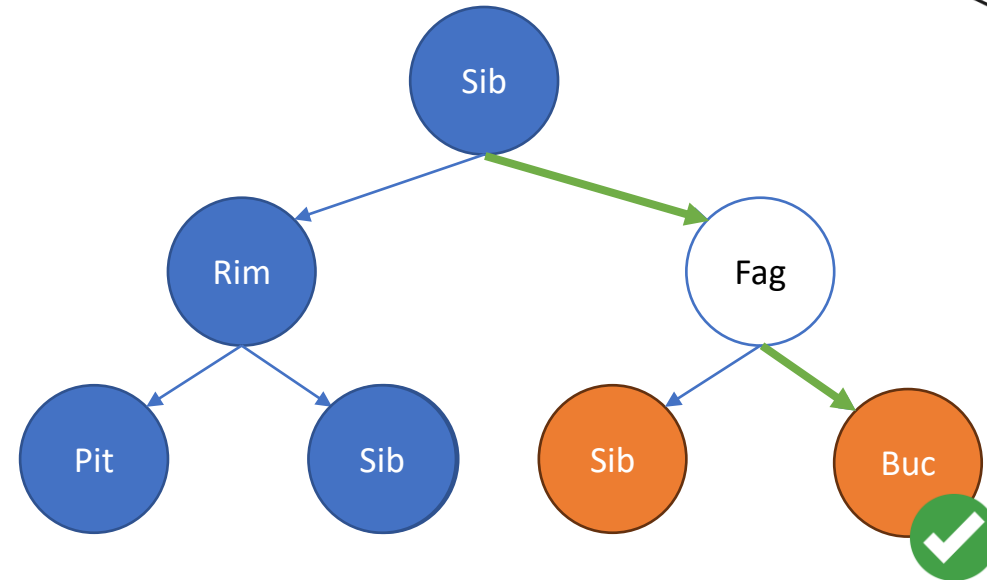
 for **action** in actions(**state**):

next state = transition(**state**, **action**)

 if **next state** is goal: return solution

frontier.add(**next state**)

return failure



Queue:



Breadth-first Search

```

create frontier : queue

insert initial state


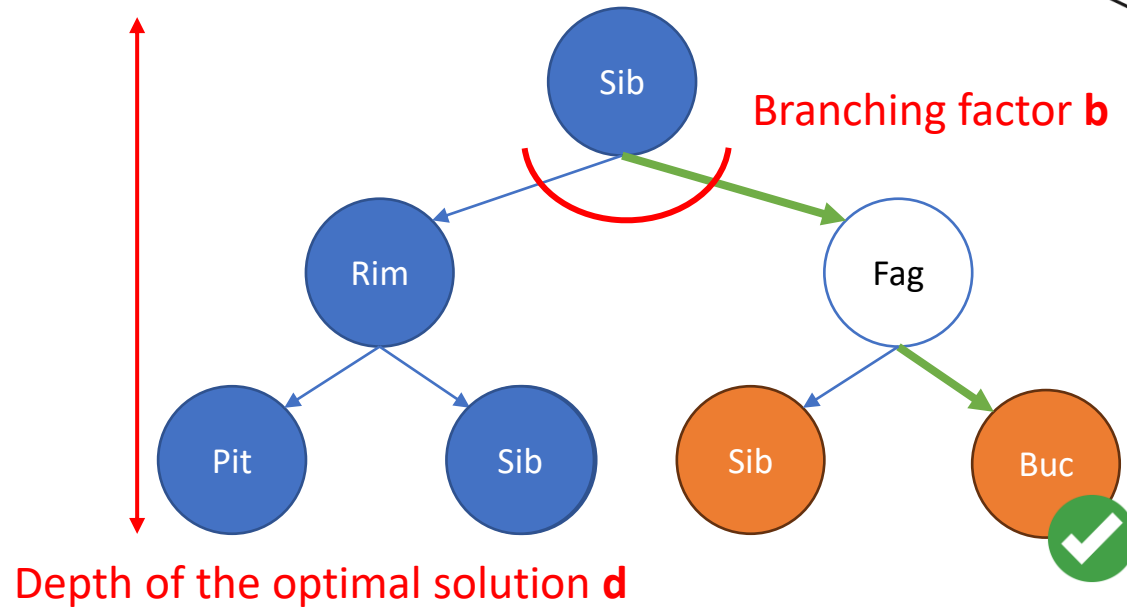
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution

        frontier.add(next state)

return failure
    
```

Queue:



- **Time complexity (# nodes expanded)?**
 - $1 + b + b^2 + \dots + b^d = O(b^d)$
- **Space complexity?**
 - $O(b^d)$, worst case: expand the last child in a branch
- **Complete?** Yes, if B is finite
- **Optimal?** Yes, if uniform cost

Uniform-cost Search

```
create frontier : priority queue (path cost)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```


```
    if state is goal: return solution
```

```
    for action in actions(state):
```

```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```



Uniform-cost Search

```
create frontier : priority queue (path cost)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```

```
    if state is goal: return solution
```



```
    for action in actions(state):
```

```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```

Priority Queue:



Uniform-cost Search

```
create frontier : priority queue (path cost)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```



```
    if state is goal: return solution
```

```
    for action in actions(state):
```


```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```



Priority Queue:



Uniform-cost Search

```
create frontier : priority queue (path cost)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```



```
    if state is goal: return solution
```

```
    for action in actions(state):
```


```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```



Priority Queue:



Uniform-cost Search

```
create frontier : priority queue (path cost)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```



```
    if state is goal: return solution
```

```
    for action in actions(state):
```


```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```



Priority Queue:



Uniform-cost Search

create **frontier** : priority queue (path cost)

insert initial state

while **frontier** is not empty:

state = **frontier**.pop()



 if **state** is goal: return solution

 for **action** in actions(**state**):


next state = transition(**state**, **action**)

frontier.add(**next state**)

return failure



Priority Queue:



Uniform-cost Search

create **frontier** : priority queue (path cost)

insert initial state

while **frontier** is not empty:

state = **frontier**.pop()



 if **state** is goal: return solution

 for **action** in actions(**state**):

next state = transition(**state**, **action**)

frontier.add(**next state**)

return failure



Priority Queue:



Uniform-cost Search

create **frontier** : **priority queue** (path cost)

insert **initial state**

while **frontier** is not empty:

state = **frontier**.pop()


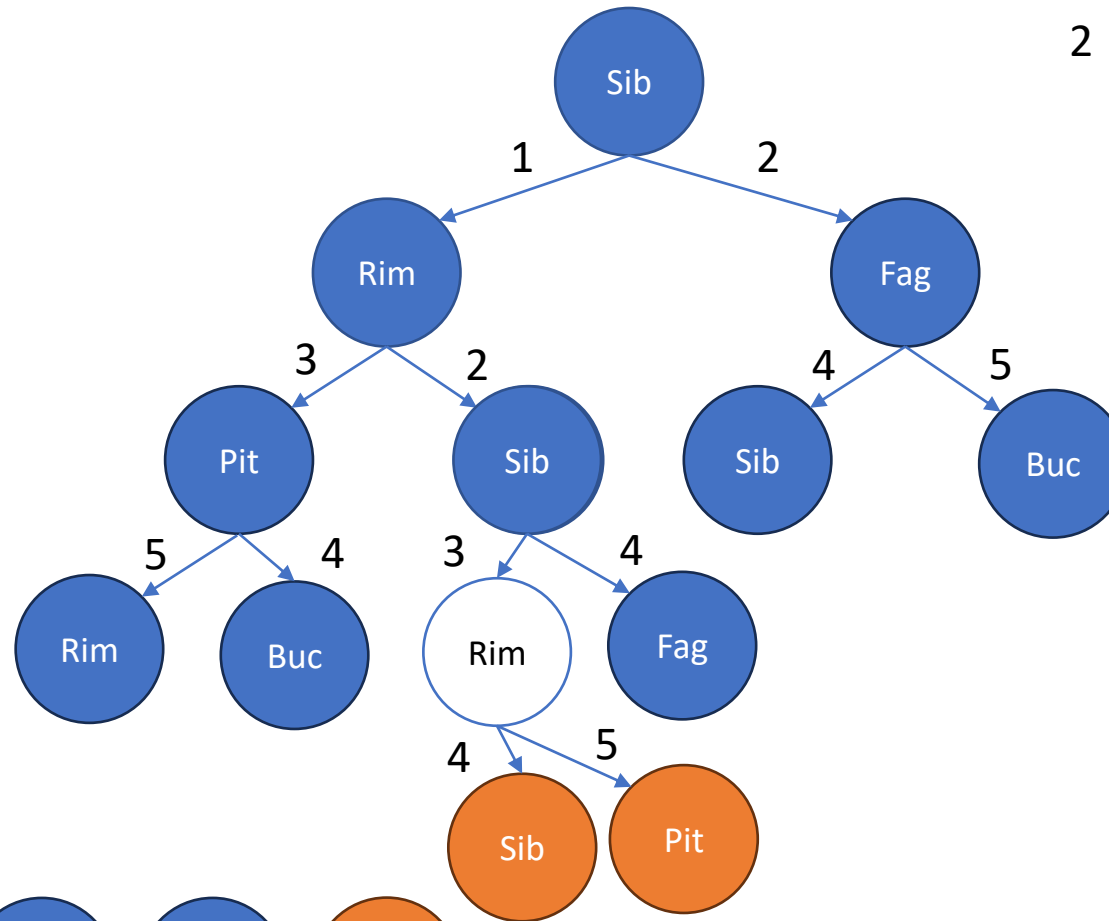
if state is goal: return solution

 for **action** in actions(**state**):


next state = transition(**state**, **action**)

frontier.add(**next state**)

return failure



Priority Queue:



Uniform-cost Search

create **frontier** : **priority queue** (path cost)

insert **initial state**

while **frontier** is not empty:

state = **frontier**.pop()


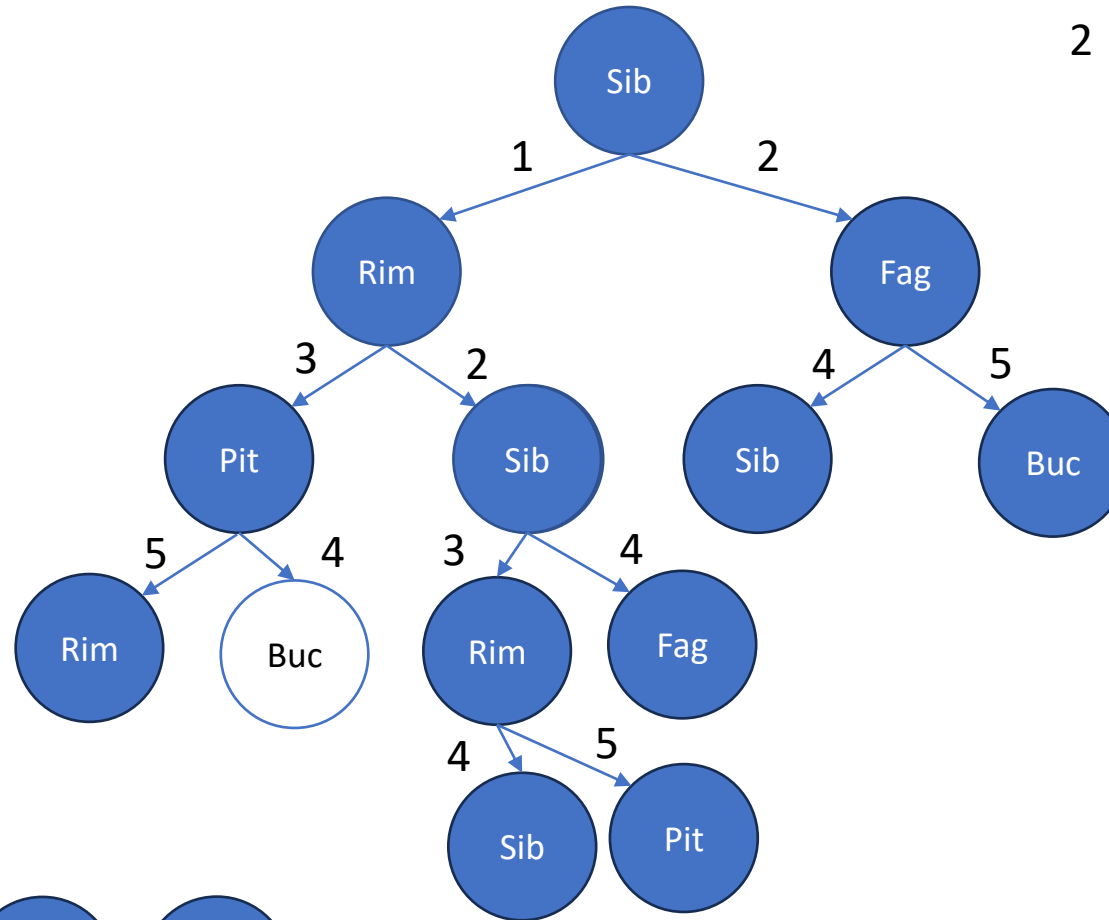
if state is goal: return solution

 for **action** in actions(**state**):


next state = transition(**state**, **action**)

frontier.add(next state)

return failure



Priority Queue:



Uniform-cost Search

create **frontier** : **priority queue** (path cost)

insert **initial state**

while **frontier** is not empty:

state = **frontier**.pop()

if state is goal: return solution




 for **action** in actions(**state**):

next state = transition(**state**, **action**)

frontier.add(next state)

return failure

Priority Queue:



- **Time complexity (# nodes expanded)?**
 - $O(b^{C^*/\epsilon})$, C^* cost of optimal solution
- **Space complexity?**
 - $O(b^{C^*/\epsilon})$
- **Complete?** Yes, if step cost $\geq \epsilon$
- **Optimal?** Yes


Depth-first Search

```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```



Depth-first Search



```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Stack:



Depth-first Search



```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Stack:



Depth-first Search




```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Stack:



Depth-first Search



```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Stack:



Depth-first Search



```
create frontier : stack

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)
return failure
```

Stack:



Depth-first Search

```
create frontier : stack

insert initial state

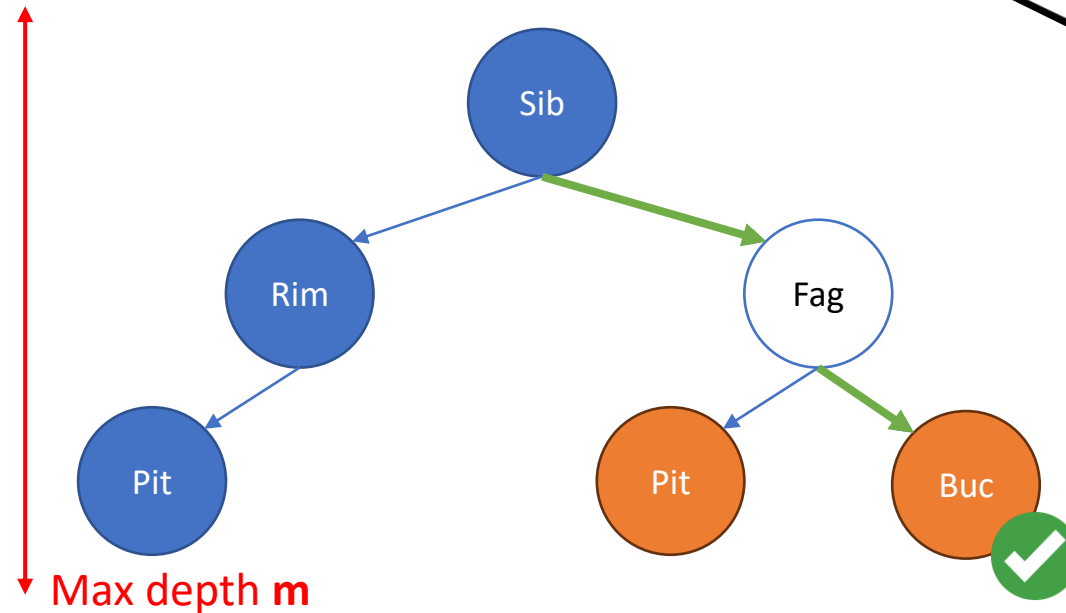
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution


        frontier.add(next state)

return failure
```

Stack:



- Time complexity (# nodes expanded)?
 - $O(b^m)$
- Space complexity?
 - $O(bm)$
- Complete?



Depth-first Search

```
create frontier : stack



insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution

    frontier.add(next state)

return failure
```



- **Time complexity (# nodes expanded)?**
 - $O(b^m)$
- **Space complexity?**
 - $O(bm)$
- **Complete?** No, when depth is infinite or can go back and forth (loops)
- **Optimal?**

Depth-first Search

```
create frontier : stack



insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution

    frontier.add(next state)

return failure
```



- **Time complexity (# nodes expanded)?**
 - $O(b^m)$
- **Space complexity?**
 - $O(bm)$
- **Complete?** No, when depth is infinite or can go back and forth (loops)
- **Optimal?** No

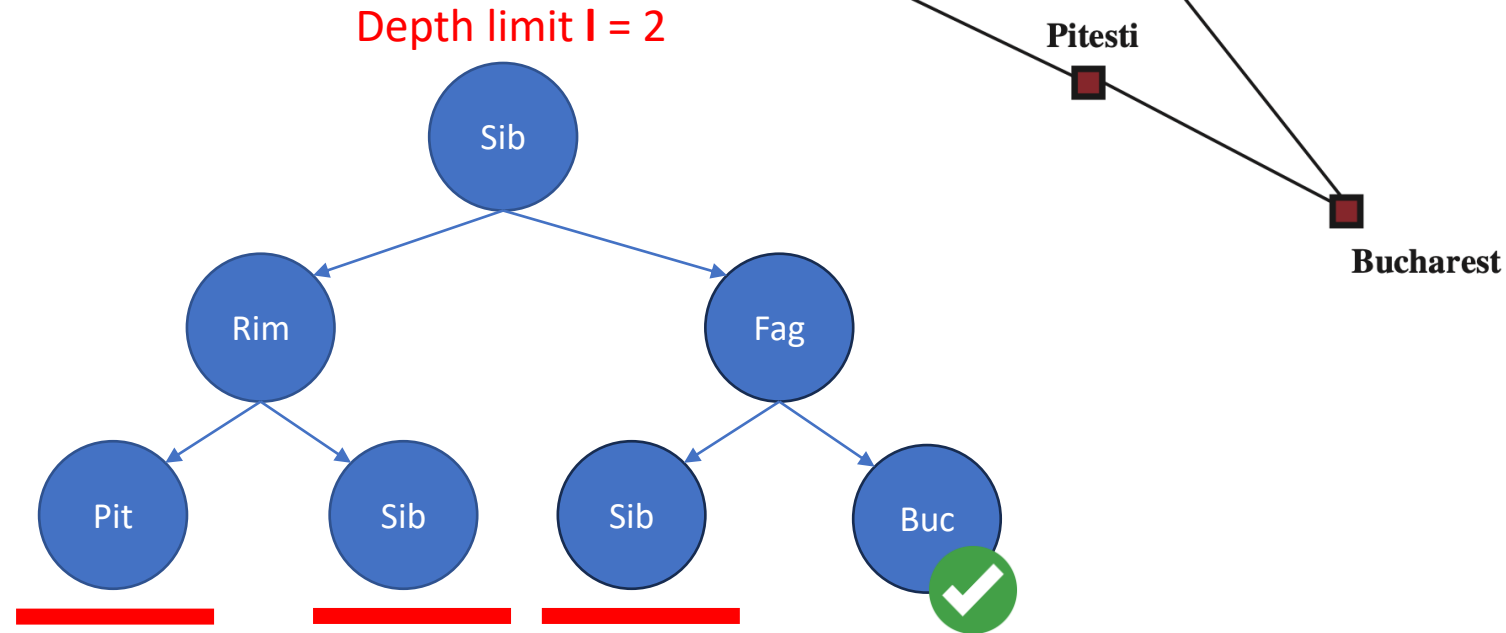
How do we handle infinite depth?

Outline

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- **Variants of uninformed search algorithms**
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Depth-limited Search (DLS)

- Limit the search depth
- Backtrack once the depth limit is reached



How do we know the depth of the solution?


We don't know :(

- Time complexity (# nodes expanded)?
 - $b^0 + b^1 + \dots + b^l = O(b^l)$
- Space complexity?
 - $O(bl)$
- Complete? No
- Optimal? No

Iterative Deepening Search (IDS)



- Do depth-limited search with max depth 0 ... N
- Return solution if found
- Increase the depth otherwise

Depth limit $l = 0$





Iterative Deepening Search (IDS)

- Do depth-limited search with max depth 0 ... N
- Return solution if found
- Increase the depth otherwise





Iterative Deepening Search (IDS)

- Do depth-limited search with max depth 0 ... N
- Return solution if found
- Increase the depth otherwise



Iterative Deepening Search (IDS)

- Do depth-limited search with max depth 0 ... N
- Return solution if found
- Increase the depth otherwise



- **Time complexity (# nodes expanded)?**
 - **Overhead!** $b^0 + (b^0 + b^1) + \dots + (b^0 + \dots + b^d)$
 - $= (d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + 2b^{d-1} + b^d = O(b^d)$
- **Space complexity?**
 - $O(bd)$
- **Complete?** Yes
- **Optimal?** Yes, if uniform step cost

Iterative Deepening Search (IDS)



For $b=10$, $d=5$,

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

Overhead = $(123,456 - 111,111)/111,111 = 11\%$

Backward Search

Search from the goal



Bidirectional Search

Combine search:

- Forward (from the start)
- Backward (from the goal)


Stop when two searches meet

Intuition:

$$2 \times O(b^{d/2}) < O(b^d)$$

Issues?

- Operators need to be reversible
- Many goal states
- How to efficiently check if a node appears in the other search tree?





Outline

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- **Dealing with repeated states**
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

Dealing with Repeated States

- Remember states that are already visited
- Don't visit again



Tree Search

Keep track of visited nodes
(use hashtable/dictionary!)

```
create frontier

insert initial state

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)

        if next state is goal: return solution
        frontier.add(next state)

return failure
```

Graph Search

```
create frontier
create visited


insert initial state to queue and visited

while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)

return failure
```

BFS with Graph Search


```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```





BFS with Graph Search

```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```

Queue:





Visited:



BFS with Graph Search

```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```



Queue:





Visited:



BFS with Graph Search

```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```



Visited!
(not added to the queue)

Queue:





Visited:



BFS with Graph Search

```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```



Queue:



Visited:




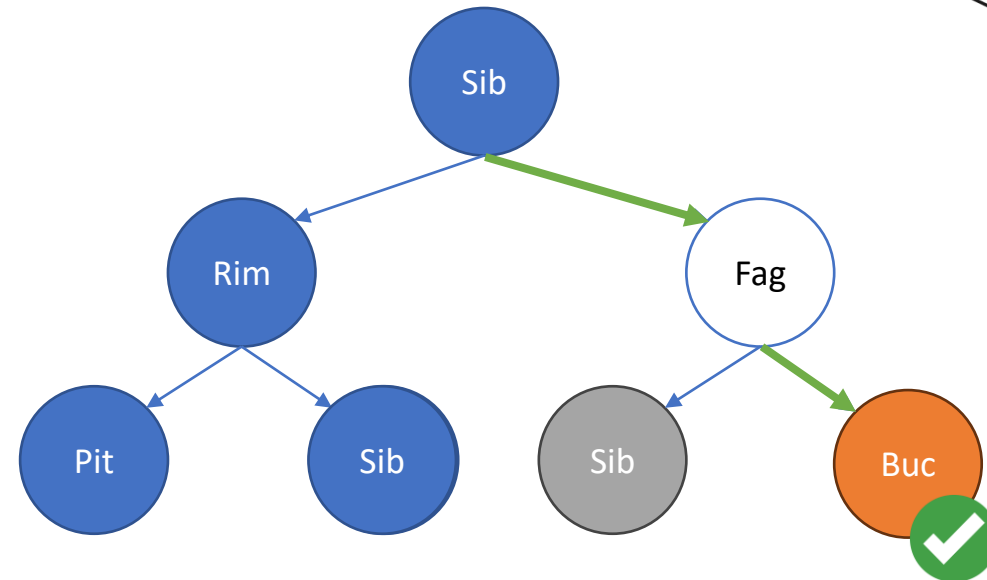
BFS with Graph Search

```
create frontier : queue
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```

Queue:




Visited:



DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```





DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```

Stack:



Visited:





DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```

Stack:




Visited:




DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```


Stack:



Visited:




Visited!
(not added to the stack)




DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```


Stack:



Visited:




Visited!
(not added to the stack)





DFS with Graph Search

```
create frontier : stack
create visited
insert initial state to queue and visited
while frontier is not empty:
    state = frontier.pop()
    for action in actions(state):
        next state = transition(state, action)
        if next state in visited: continue
        if next state is goal: return solution
        frontier.add(next state)
        visited.add(next state)
return failure
```

Stack:



Visited:



Choosing a Search Strategy

Depends on the problem:

- Number of goal states
- Distribution of goal states in search tree
- Finite/infinite branching factor/depth
- Repeated states
- Need for optimality?
- Need to know if there is no solution?

Summary: Uninformed Search Algorithms

- Search algorithms
 - Breadth-first search – **queue**, explore layer by layer
 - Uniform-cost search – **priority queue** (path cost)
 - Depth-first search – **stack**, go deep first then backtrack
- Variants:
 - Depth limited search – **limit max depth** of the search
 - Iterative deepening search – try DLS with depth limit 0, ..., N
 - Bidirectional search – search from the **start** and the **goal**, meet in the **middle**
- Dealing with repeated states
 - Graph search – don't visit nodes that are already **visited**

Outline

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- **Informed search algorithms**
 - Greedy best-first search
 - A* search
 - Heuristics
- Variants of A*

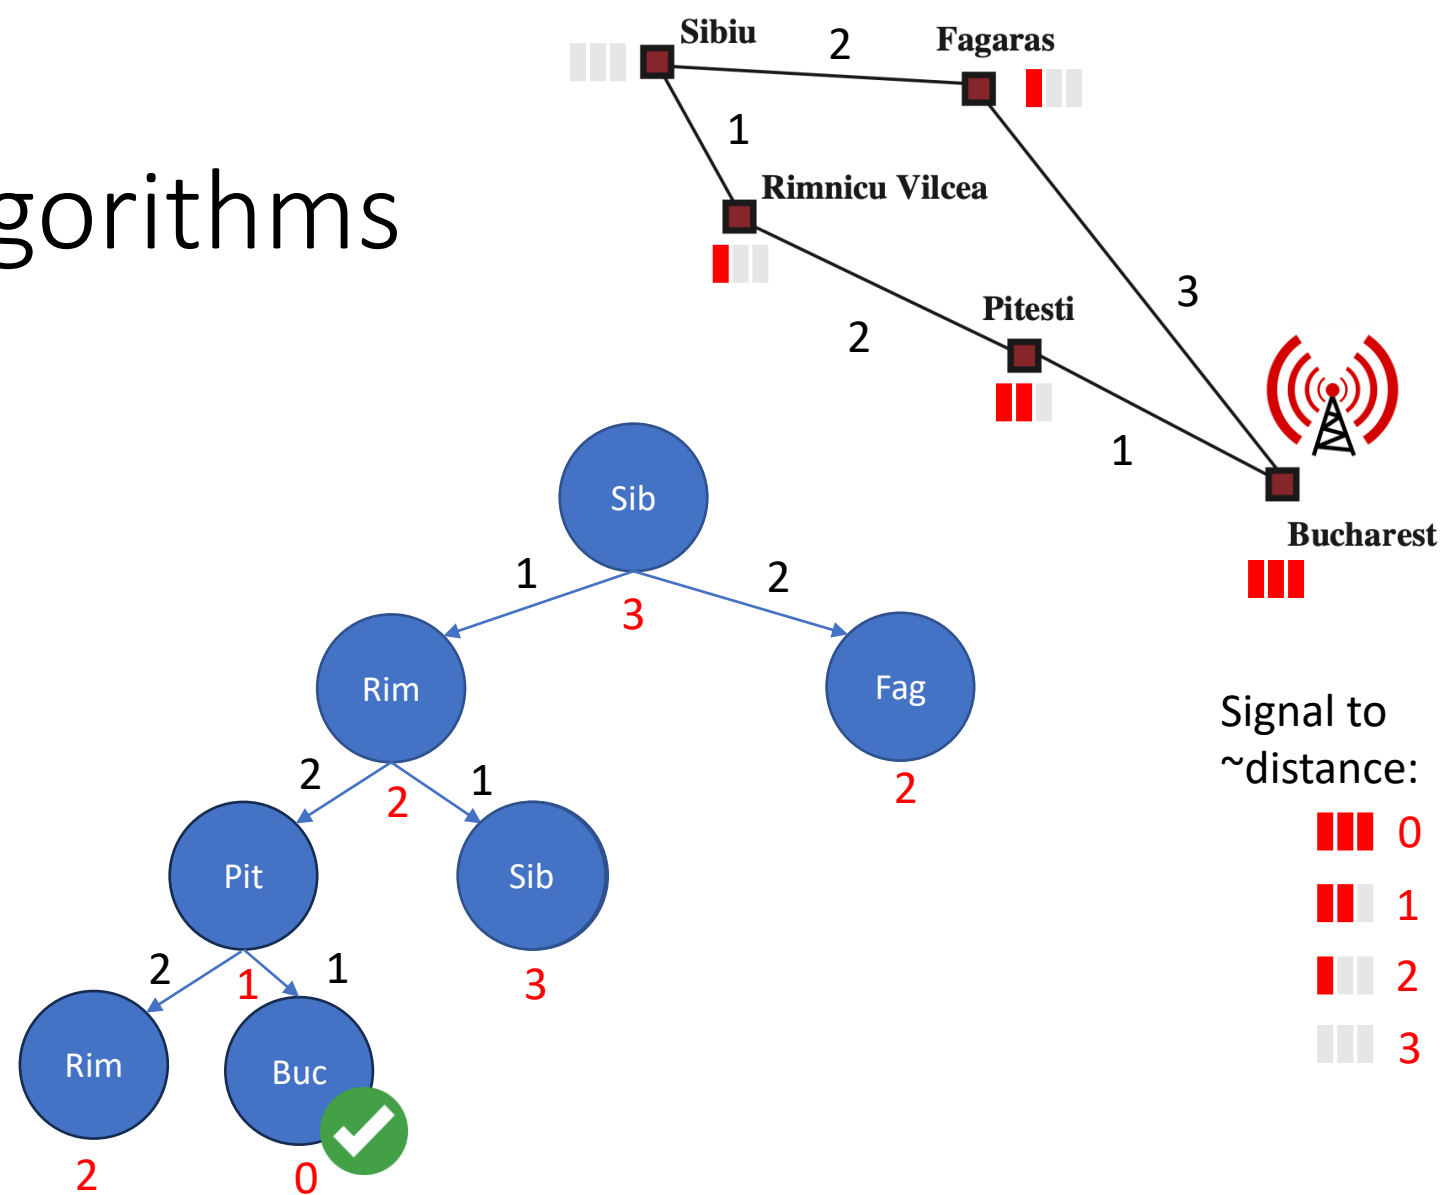
Informed Search Algorithms

Uninformed search:

Search blindly

Informed search:

Use domain information to guide the search



Best-first Search

```
create frontier : priority queue f(n)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```

```
    if state is goal: return solution
```



```
    for action in actions(state):
```

```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```

Evaluation function $f(n)$:
estimate the "goodness" of a state



Signal to
~distance:



Special cases:

- Greedy best-first search
- A* search

Greedy Best-first Search

create **frontier** : priority queue **f(n)**

insert initial state

while **frontier** is not empty:

state = **frontier**.pop()

 if **state** is goal: return solution

 for **action** in actions(**state**):

next state = transition(**state**, **action**)



frontier.add(**next state**)

return failure

Evaluation function f(n):
estimate the "goodness" of a state

f(n) = h(n)

Heuristic: estimated cost from n to goal



Signal to
~distance:

■■■■ 0
■■■ 1
■■ 2
■ 3

- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete?**

Greedy Best-first Search

create **frontier** : priority queue $f(n)$

insert initial state

while **frontier** is not empty:

$state = frontier.pop()$

 if $state$ is goal: return solution

 for $action$ in actions($state$):

$next\ state = transition(state, action)$

$frontier.add(next\ state)$

return failure

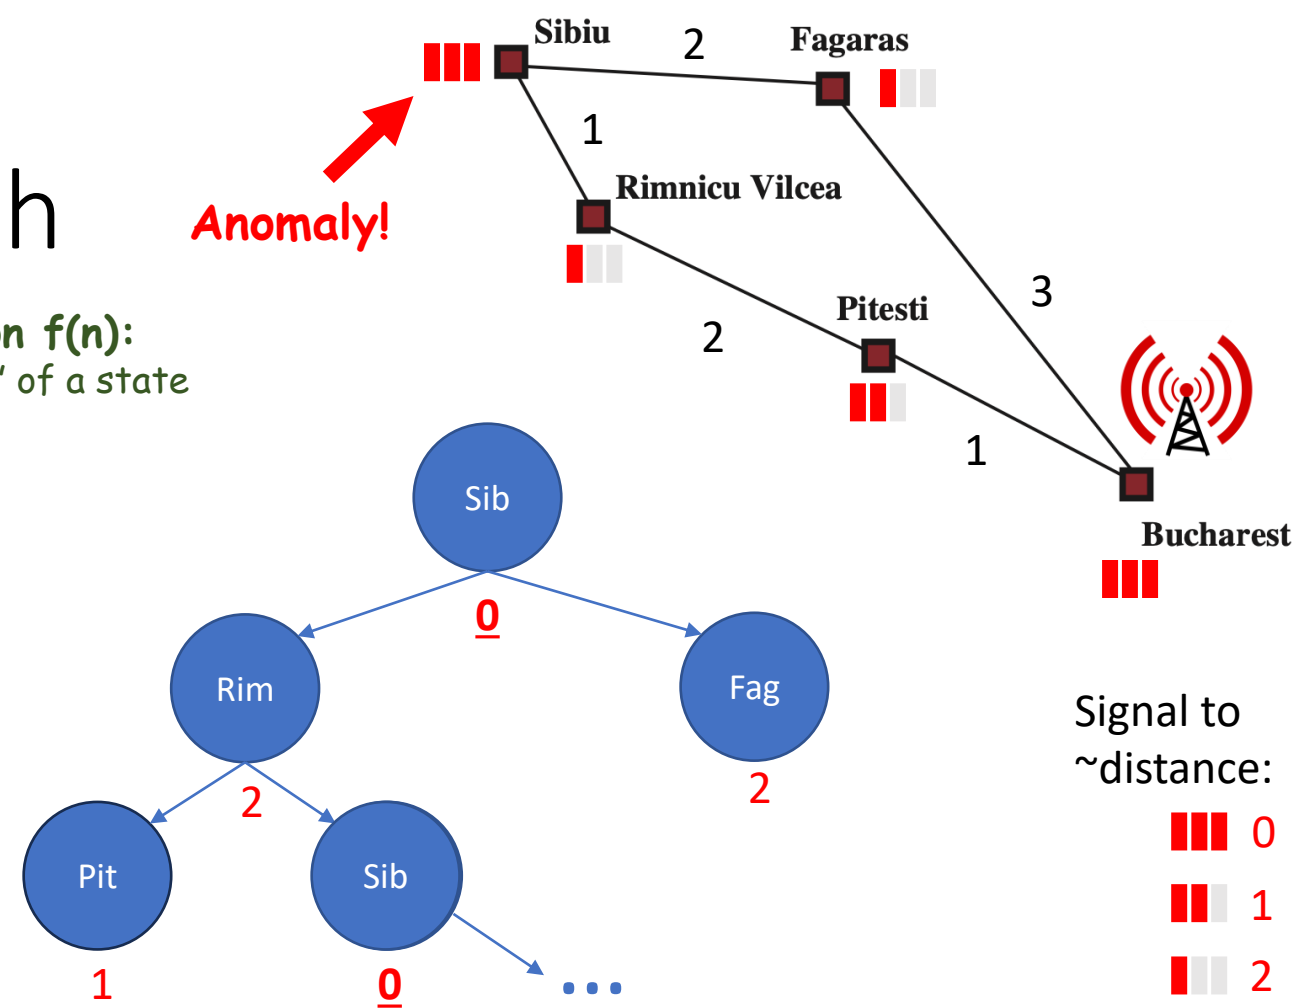
Evaluation function $f(n)$:
estimate the "goodness" of a state

$f(n) = h(n)$

Heuristic: estimated cost from n to goal

Doesn't consider the cost so far!

Anomaly!



- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete? No**
- **Optimal? No**

A* Search

create **frontier** : priority queue $f(n)$

insert initial state

while **frontier** is not empty:

$state = frontier.pop()$

 if $state$ is goal: return solution

 for $action$ in $actions(state)$:

$next\ state = transition(state, action)$

$frontier.add(next\ state)$

return failure

Evaluation function $f(n)$:
estimate the "goodness" of a state



$$f(n) = g(n) + h(n)$$



Cost so far to reach n



Heuristic: estimated cost from n to goal



Signal to
~distance:




- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete?** Yes
- **Optimal?** Yes*

Admissible Heuristics

A **heuristic $h(n)$** is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true cost** to reach the goal state from n .

An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is a **conservative estimate**.

Theorem: if $h(n)$ is admissible, A^* using **tree search** is optimal



Example: $h_{SLD}(n)$ never overestimates the actual road distance


Admissible Heuristics: Optimality

Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .

$$\begin{aligned}f(G) &= g(G) + h(G) = g(G) + 0 \\f(G_2) &= g(G_2) + h(G_2) = g(G_2) + 0 \\g(G_2) &> g(G) \text{ since } G_2 \text{ is suboptimal} \\f(G_2) &> f(G)\end{aligned}$$

$$\begin{aligned}h(n) &\leq h^*(n) \text{ since } h \text{ is admissible} \\f(n) &= g(n) + h(n) \leq g(n) + h^*(n) = f(G)\end{aligned}$$

$$f(n) \leq f(G) < f(G_2), \text{ } G_2 \text{ will never be expanded}$$




Consistent Heuristics

A **heuristic $h(n)$** is **consistent** if for every node n , every successor n' of n generated by any action a , $h(n) \leq c(n,a,n') + h(n')$

If h is consistent, we have


$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$

i.e., $f(n)$ is **non-decreasing** along any path



Theorem: If $h(n)$ is consistent, A^* using **graph search** is optimal

Consistent Heuristics: Optimality



A * expands nodes in order of increasing f-cost

Gradually adds "f-contours" of nodes

Contour i has all nodes with $f = f^{(i)}$, where $f^{(i)} < f^{(i+1)}$

Dominance

If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 dominates h_1 .
 h_2 is better for search.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Heuristics

- h_1 number of misplaced tiles
- h_2 total Manhattan distance

h_2 dominates h_1

If each tile is at most one distance away from the goal, then $h_2 = h_1$, otherwise $h_2 > h_1$

“Inventing” Admissible Heuristics

A problem with **fewer restrictions** on the actions is called a relaxed problem. The cost of an **optimal solution** to a relaxed problem is an **admissible heuristic** for the original problem.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Original:

A tile can only move to adjacent blank

Relaxations:

- Each tile can move anywhere
 - h_1 number of misplaced tiles
- Each tile can move to any adjacent square
 - h_2 total Manhattan distance

Outline

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states
- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
- **Variants of A***

Variants of A*

- Iterative Deepening A* (IDA*)
 - Use iterative deepening search
 - Cutoff using f-cost [$f(n) = g(n) + h(n)$] instead of depth
- Simplified Memory-bounded A* (SMA*)
 - Drop the nodes with worst f-cost if memory is full

Summary: Informed Search Algorithms

- Informed search: guide search with domain information
- Best-first search
 - Greedy best-first search
 - $f(n) = h(n)$, **heuristic** estimate of cost from n to goal
 - A* search
 - $f(n) = g(n) + h(n)$, cost so far + heuristic
- Heuristics
 - Admissible: $h(n) \leq h^*(n)$
 - Consistent: $h(n) \leq c(n, a, n') + h(n')$
 - Dominant: if $h_1(n) \leq h_2(n)$, h_2 dominant
- Creating admissible heuristic: **true cost** of the **relaxed problem**
- A* variants: IDA*, SMA*
 - Idea: prune to save memory

Summary

- Problem-solving agents
- Uninformed search algorithms
 - Breadth-first Search (BFS) – layer by layer
 - Uniform-cost search – Dijkstra
 - Depth-first Search (DFS) – go deep first
- Variants of uninformed search algorithms
 - Depth-limited search – set max depth
 - Iterative deepening search – try DLS with depth limit 0, ..., N
 - Bidirectional search – combine forward and backward search
- Dealing with repeated states – visit state only once
- Informed search algorithms
 - Greedy best-first search – $f(n) = h(n)$
 - A* search – $f(n) = g(n) + h(n)$
 - Heuristics: admissibility, consistency, dominance
- Variants of A*: IDA*, SMA*

Coming Up Next Week

- Local search
 - Hill climbing
 - Simulated annealing
 - Beam search
 - Genetic algorithms
- Adversarial search
 - Games vs search problems
 - Minimax
 - Alpha-beta pruning

To Do

- **Lecture Training 2**
 - +100 Free EXP
 - +50 Early bird bonus
- **Problem Set 0**
 - Due Saturday, 26th August (Tomorrow)
- **Tutorial Swaps**
 - Due Sunday, 27th August