

CS2109S: Introduction to AI and Machine Learning

Lecture 3: Informed, Local, and Adversarial Search

25 August 2023

Recap

- Problem-solving agents
- Search algorithms
- Uninformed search algorithms
 - Breadth-first Search (BFS)
 - Uniform-cost search
 - Depth-first Search (DFS)
- Variants of uninformed search algorithms
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search
- Dealing with repeated states

Questions from Last Week

Uniform-cost search

- **Why the name?**
 - Exploring states of the same path cost uniformly
- **Why $O(b^{C^*}/\epsilon)$ is the upperbound depth although there are other nodes explored?**
 - The depth of the other nodes explored are less than the bound
- **Why the condition that step cost $\geq \epsilon$?**
 - This is not a condition, but simply an assumption for analysis

Outline

- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
 - Variants of A*
- Local search
 - Hill climbing
 - Simulated annealing
 - Beam search
 - Genetic algorithms
- Adversarial search
 - Games
 - Minimax
 - Alpha-beta pruning
 - Handling large/infinite game trees
 - Optimizing the search

} Recap!

Outline

- **Informed search algorithms**

- Greedy best-first search
- A* search
- Heuristics
- Variants of A*

} Recap!

- Local search

- Hill climbing
- Simulated annealing
- Beam search
- Genetic algorithms

- Adversarial search

- Games
- Minimax
- Alpha-beta pruning
- Handling large/infinite game trees
- Optimizing the search

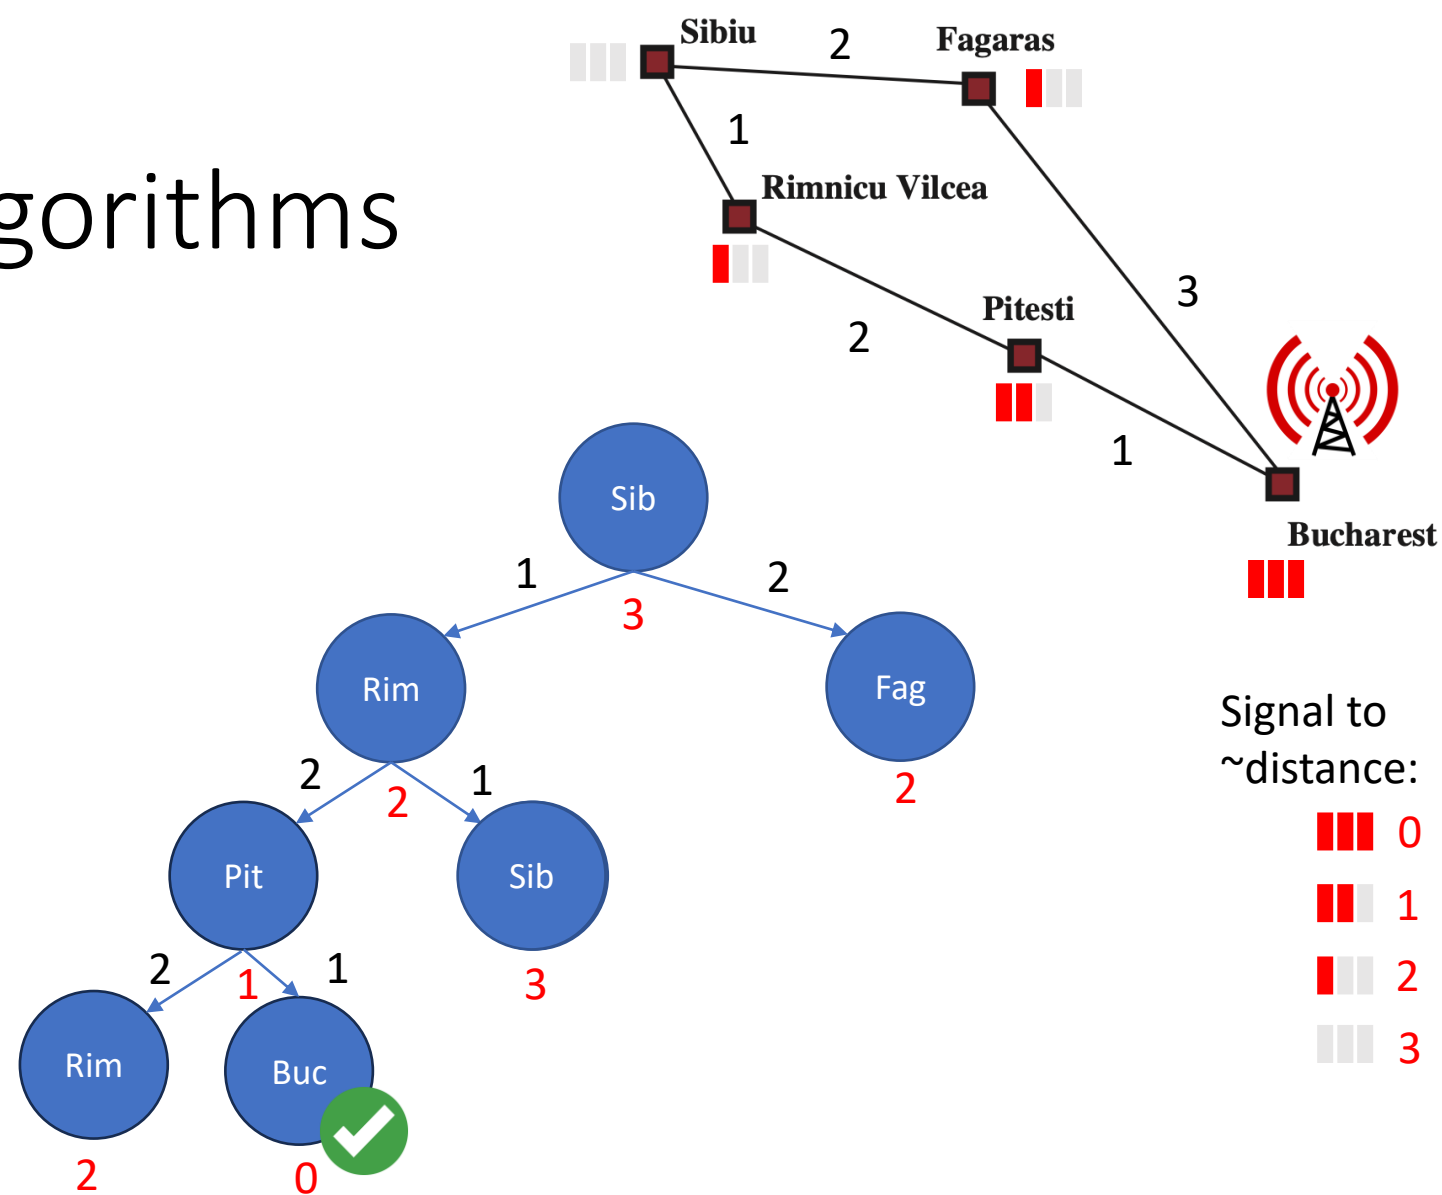
Informed Search Algorithms

Uninformed search:

Search blindly

Informed search:

Use domain information to guide the search



Best-first Search

```
create frontier : priority queue f(n)
```

```
insert initial state
```

```
while frontier is not empty:
```

```
    state = frontier.pop()
```

```
    if state is goal: return solution
```



```
    for action in actions(state):
```

```
        next state = transition(state, action)
```

```
        frontier.add(next state)
```

```
return failure
```

Evaluation function $f(n)$:
estimate the "goodness" of a state



Signal to
~distance:



Special cases:

- Greedy best-first search
- A* search

Greedy Best-first Search

create **frontier** : priority queue **f(n)**

insert **initial state**

while **frontier** is not empty:

state = **frontier**.pop()

 if **state** is goal: return solution

 for **action** in actions(**state**):

next state = transition(**state**, **action**)



frontier.add(**next state**)

return failure

Evaluation function f(n):
estimate the "goodness" of a state

f(n) = h(n)

Heuristic: estimated cost from n to goal



Signal to
~distance:

■■■ 0
■■ 1
■ 2
 3

- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete?**

Greedy Best-first Search

create **frontier** : priority queue $f(n)$

insert initial state

while **frontier** is not empty:

$state = frontier.pop()$

 if $state$ is goal: return solution

 for $action$ in actions($state$):

$next\ state = transition(state, action)$

$frontier.add(next\ state)$

return failure

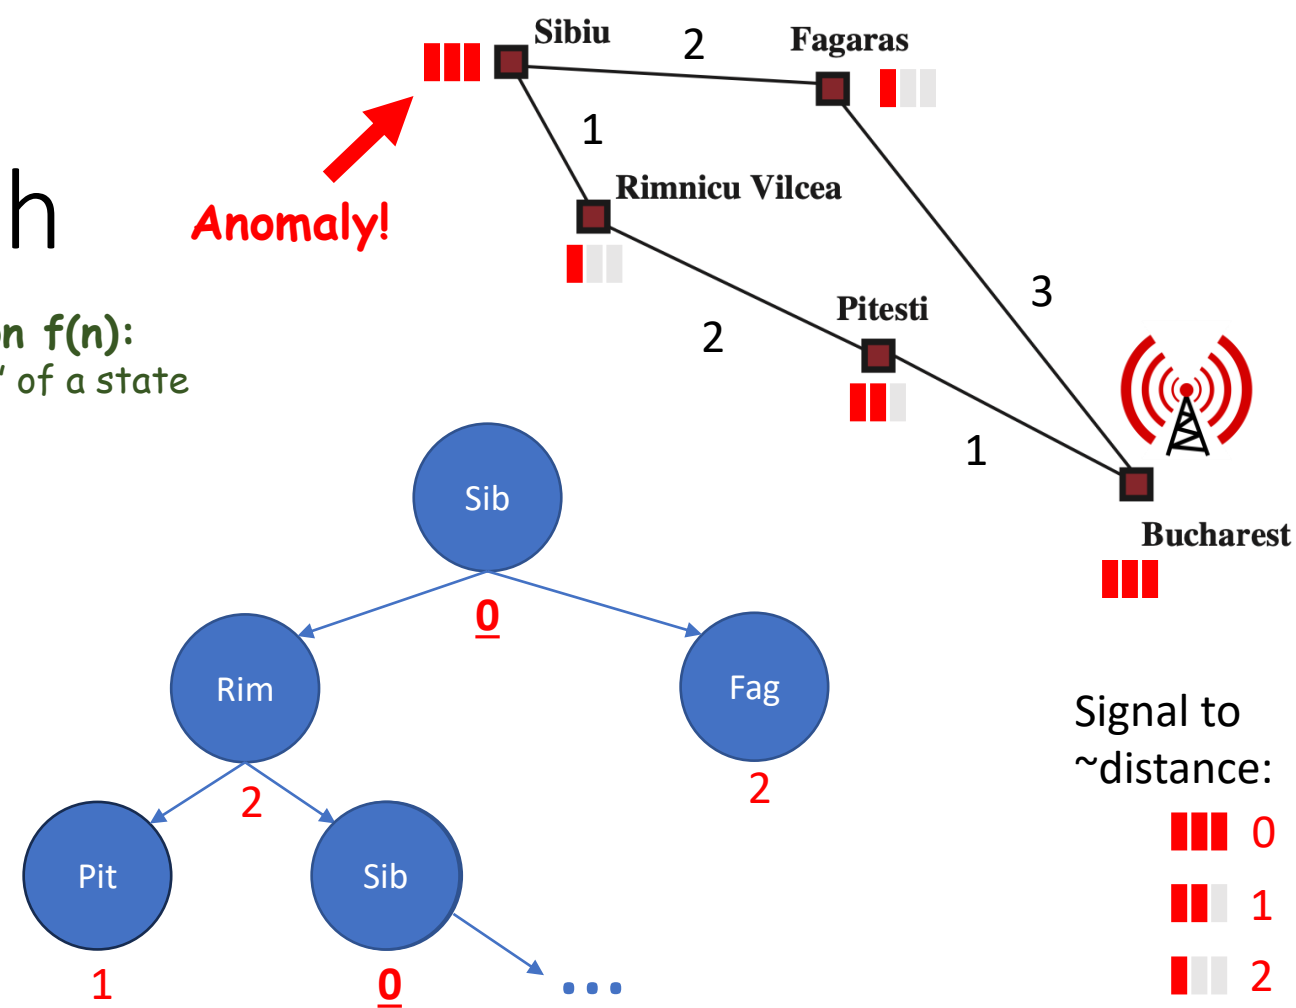
Evaluation function $f(n)$:
estimate the "goodness" of a state

$f(n) = h(n)$

Heuristic: estimated cost from n to goal

Doesn't consider the cost so far!

Anomaly!



- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete? No**
- **Optimal? No**

A* Search

create **frontier** : priority queue $f(n)$

insert initial state

while **frontier** is not empty:

$state = frontier.pop()$

 if $state$ is goal: return solution

 for $action$ in $actions(state)$:

$next\ state = transition(state, action)$

$frontier.add(next\ state)$

return failure

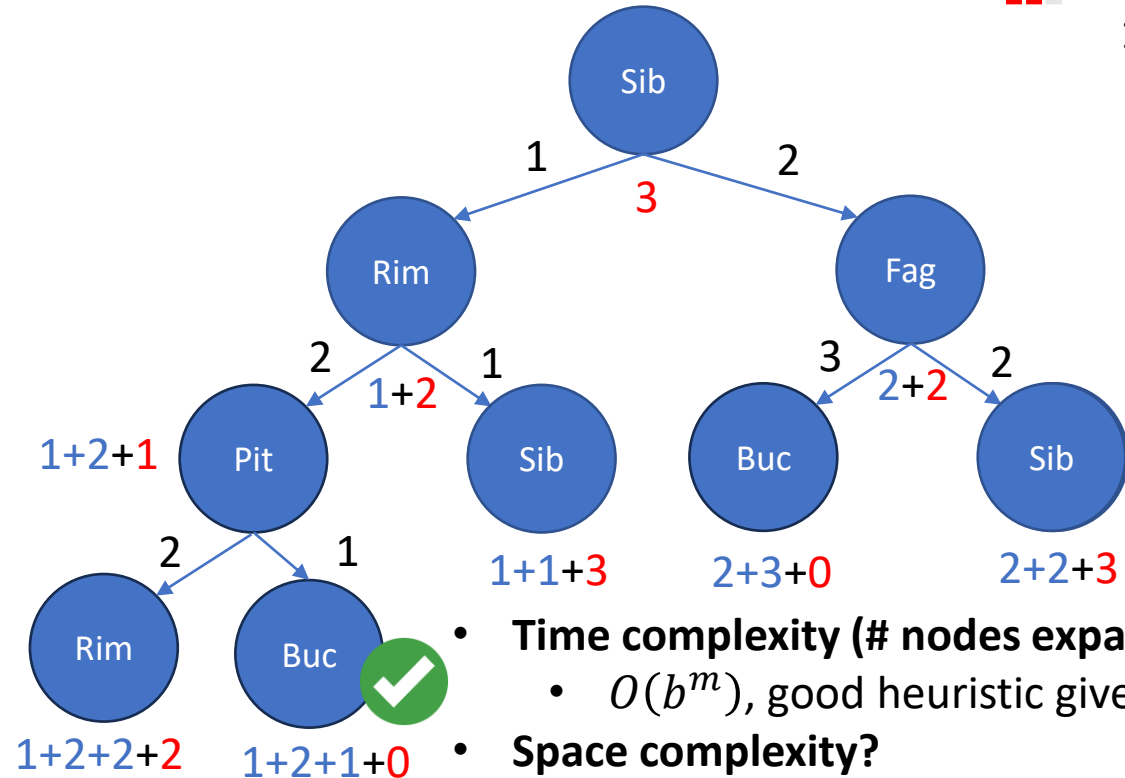

Evaluation function $f(n)$:
estimate the "goodness" of a state

$$f(n) = g(n) + h(n)$$



Cost so far to reach n

Heuristic: estimated cost from n to goal




- **Time complexity (# nodes expanded)?**
 - $O(b^m)$, good heuristic gives improvement
- **Space complexity?**
 - $O(b^m)$, keep all nodes in memory
- **Complete?** Yes
- **Optimal?** Yes*

Admissible Heuristics

A **heuristic $h(n)$** is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true cost** to reach the goal state from n .

An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is a **conservative estimate**.

Theorem: if $h(n)$ is admissible, A^* using **tree search** is optimal



Example: $h_{SLD}(n)$ never overestimates the actual road distance


Admissible Heuristics: Optimality

Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .

$$\begin{aligned}f(G) &= g(G) + h(G) = g(G) + 0 \\f(G_2) &= g(G_2) + h(G_2) = g(G_2) + 0 \\g(G_2) &> g(G) \text{ since } G_2 \text{ is suboptimal} \\f(G_2) &> f(G)\end{aligned}$$

$$\begin{aligned}h(n) &\leq h^*(n) \text{ since } h \text{ is admissible} \\f(n) &= g(n) + h(n) \leq g(n) + h^*(n) = f(G)\end{aligned}$$

$$f(n) \leq f(G) < f(G_2), \text{ } G_2 \text{ will never be expanded}$$




Consistent Heuristics

A **heuristic $h(n)$** is **consistent** if for every node n , every successor n' of n generated by any action a , $h(n) \leq c(n,a,n') + h(n')$

If h is consistent, we have


$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$

i.e., $f(n)$ is **non-decreasing** along any path



Theorem: If $h(n)$ is consistent, A^* using **graph search** is optimal

Consistent Heuristics: Optimality



A * expands nodes in order of increasing f-cost

Gradually adds "f-contours" of nodes

Contour i has all nodes with $f = f^{(i)}$, where $f^{(i)} < f^{(i+1)}$

Dominance

If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 dominates h_1 .
 h_2 is better for search.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Heuristics

- h_1 number of misplaced tiles
- h_2 total Manhattan distance

h_2 dominates h_1

If each tile is at most one distance away from the goal, then $h_2 = h_1$, otherwise $h_2 > h_1$

“Inventing” Admissible Heuristics

A problem with **fewer restrictions** on the actions is called a relaxed problem. The cost of an **optimal solution** to a relaxed problem is an **admissible heuristic** for the original problem.

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Original:

A tile can only move to adjacent blank

Relaxations:

- Each tile can move anywhere
 - h_1 number of misplaced tiles
- Each tile can move to any adjacent square
 - h_2 total Manhattan distance

Variants of A*

- Iterative Deepening A* (IDA*)
 - Use iterative deepening search
 - Cutoff using f-cost [$f(n) = g(n) + h(n)$] instead of depth
- Simplified Memory-bounded A* (SMA*)
 - Drop the nodes with worst f-cost if memory is full

Summary: Informed Search

- Informed search: guide search with domain information
- Best-first search
 - Greedy best-first search
 - $f(n) = h(n)$, **heuristic** estimate of cost from n to goal
 - A* search
 - $f(n) = g(n) + h(n)$, cost so far + heuristic
- Heuristics
 - Admissible: $h(n) \leq h^*(n)$
 - Consistent: $h(n) \leq c(n, a, n') + h(n')$
 - Dominant: if $h_1(n) \leq h_2(n)$, h_2 dominant
- Creating admissible heuristic: **true cost** of the **relaxed problem**
- A* variants: IDA*, SMA*
 - Idea: prune to save memory


Outline

- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
 - Variants of A*
- **Local search**
 - Hill climbing
 - Simulated annealing
 - Beam search
 - Genetic algorithms
- Adversarial search
 - Games
 - Minimax
 - Alpha-beta pruning
 - Handling large/infinite game trees
 - Optimizing the search

Local Search


Previously: path to a goal is a solution

There are problems where path is not important (state is the solution).



3			8	1			2
2		1		3		6	4
			2	4			
8		9				1	6
	6						5
7		2				4	9
			5	9			
9		4		8		7	5
6			1	7			3

Credit: Encyclopaedia Britannica




Credit: Frontiersin


How do we solve these problems? **Start with a state** (solution) then **try to improve it**

Trivial Algorithms

- Random sampling
 - Generate a state randomly until a solution is found



- Random walk
 - Randomly pick a neighbour from the current state




Can we do better?

Formulation

Formally, a local search problem can be formulated as follows.

- **States (state space)**
 - State representation: grid with N queens
- **Initial state**
 - N queens positioned randomly on the grid
- **Goal test**
 - No queens are attacking each other
- **Successor function**
 - Move a single queen to another square on the same column




Evaluation Function

Output the value (“goodness”) of a state.

Can be either:

- Heuristic function
 - # of well-positioned queens
- Objective function
 - $J(x) = -x^2$



Hill climbing algorithm

current = initial state



loop:

neighbor = a **highest value** successor of **current**

if **value**(**neighbor**) <= **value**(**current**):

return **current**

current = **neighbor**




Simulated Annealing

```
current = initial state
for t = 1 ... ∞:
    T = schedule(t)
    if T = 0:
        return current
    next = a randomly selected successor of current
    if value(next) > value(current) or Prob(next, current, T):
        current = next
```

$$\text{Prob}(\text{next}, \text{current}, T) = e^{\frac{\text{value}(\text{next}) - \text{value}(\text{current})}{T}}$$

Allow “bad moves”
from time to time

Select a “bad move”



Theorem: if T decreases slowly enough, simulated annealing will find a global optimum with high probability


Beam Search

Perform k hill-climbing in parallel

Variants:

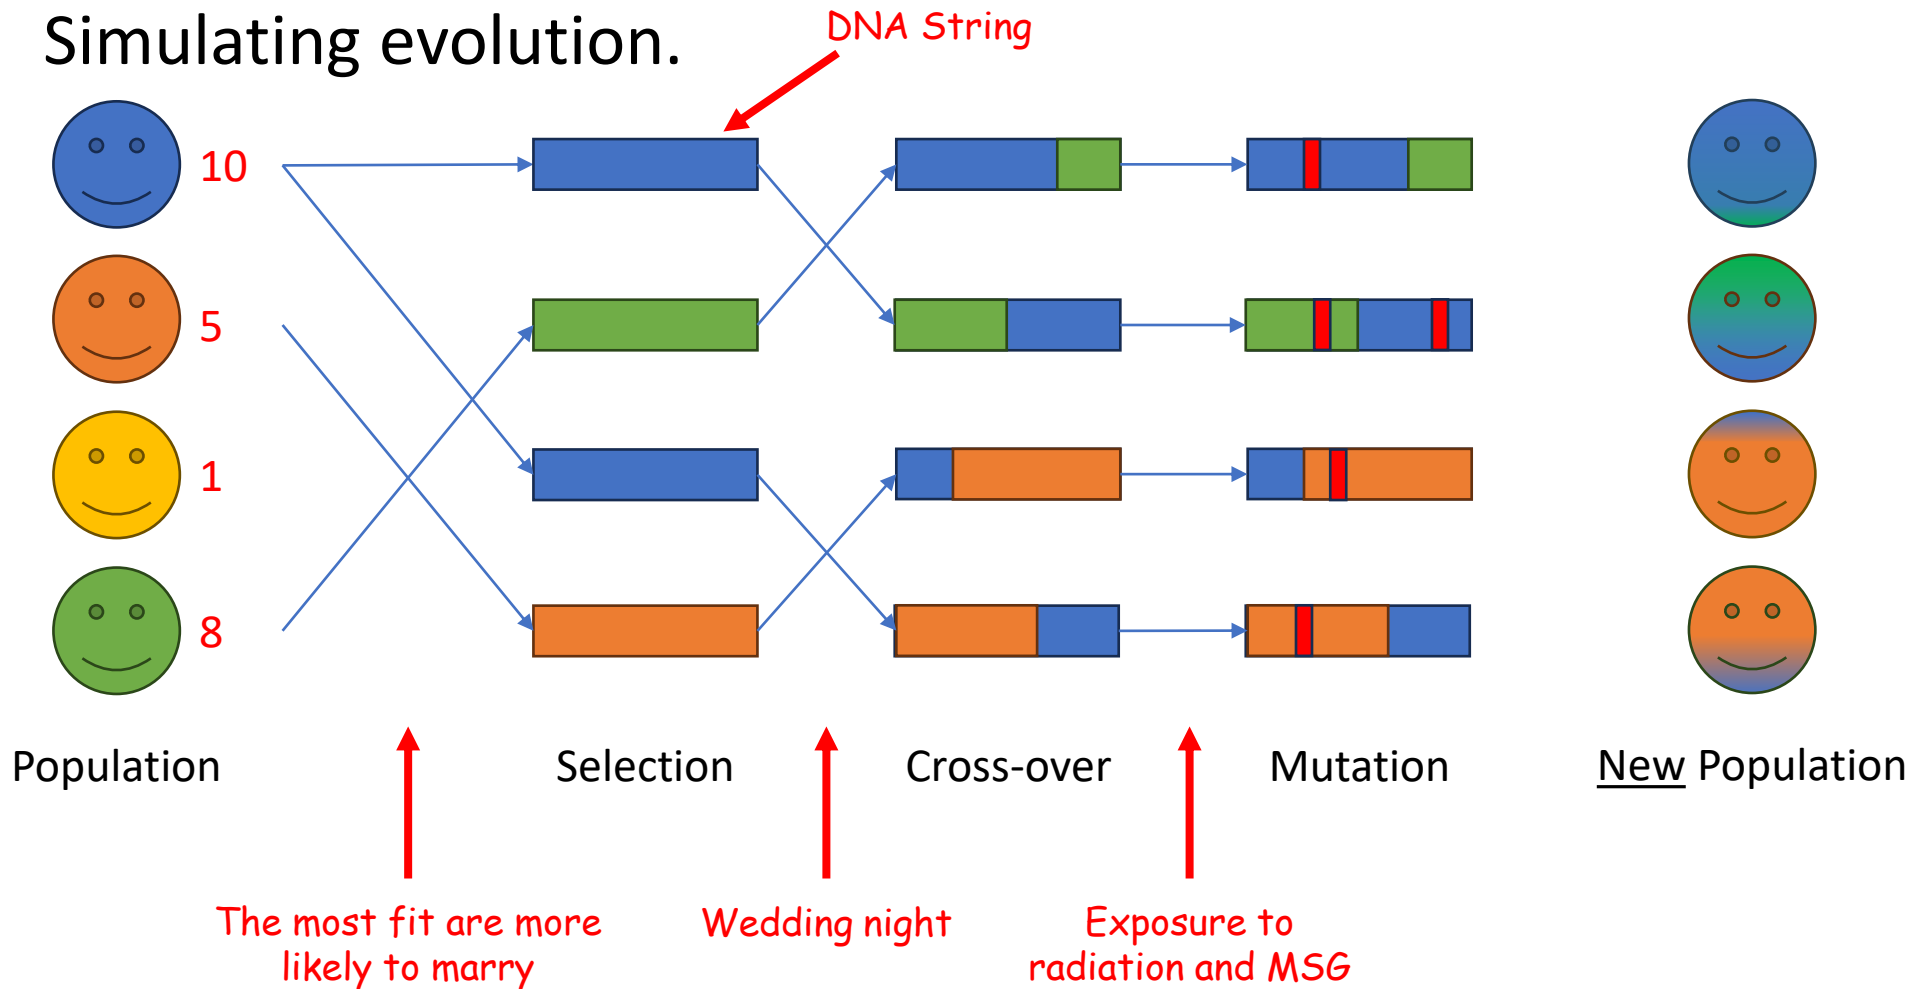
- Local beam search
 - Choose k successors deterministically
- Stochastic beam search
 - Choose k successors probabilistically

Beam width $k = 2$




Genetic Algorithm


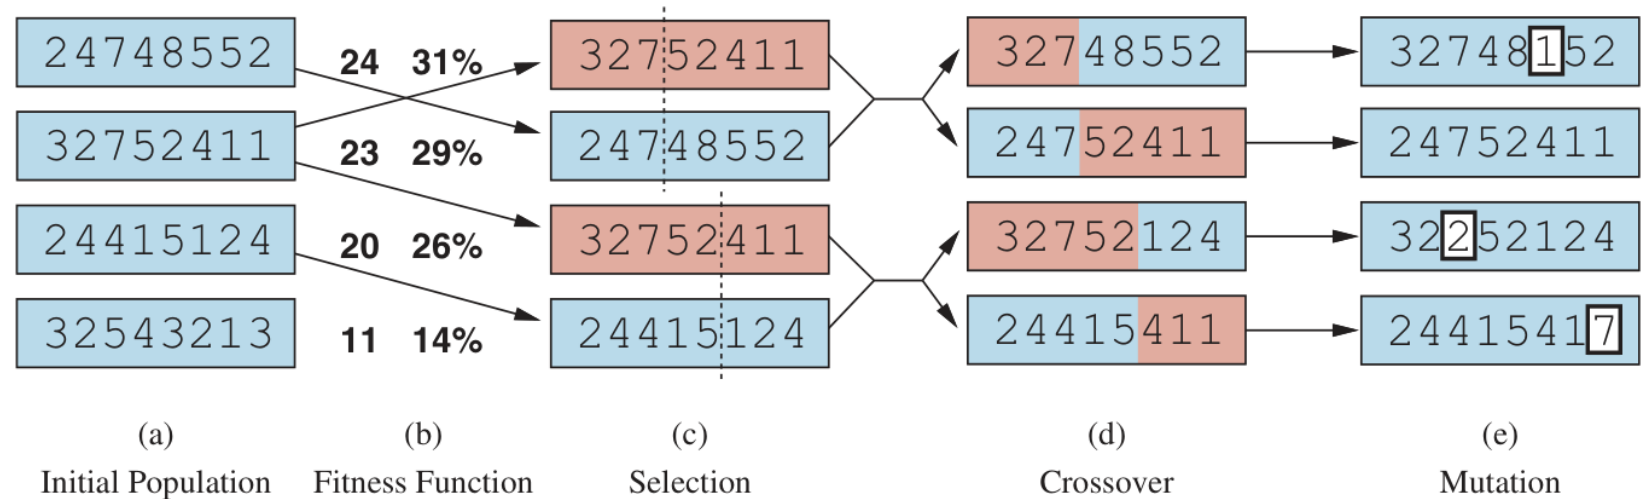
Simulating evolution.



Genetic Algorithm: 8-Queens



DNA String: 16257483



Summary: Local Search

- Local search: **path is not important**, just want the state
 - Goal state is unknown
- Trivial algorithms
 - Random sampling: **random sample a state** until solution is found
 - Random walk: **go to random neighbours** until solution is found
- Algorithms:
 - Hill-climbing: **pick the best** among neighbours, repeat
 - Simulated Annealing: hill-climbing but **allow bad moves** from time to time
 - Beam search: do **k hill-climbing in parallel**
 - Genetic algorithm: **marry** the best, **mutate**, repeat


Outline

- Informed search algorithms
 - Greedy best-first search
 - A* search
 - Heuristics
 - Variants of A*
- Local search
 - Hill climbing
 - Simulated annealing
 - Beam search
 - Genetic algorithms
- **Adversarial search**
 - Games
 - Minimax
 - Alpha-beta pruning
 - Handling large/infinite game trees
 - Optimizing the search

Games

Search problems:

The environment reacts either **deterministically** or **stochastically**




Games:

- The opponent reacts “unpredictably” (**strategic** environment)
- Usually have time limits




Example: Tic-tac-toe



Example: Game of NIM


- Several piles of sticks are given
 - Representation: **monotone** sequence of integers, e.g., (1,3,5)
- A player may **remove any number of sticks from one pile** in one turn
 - Example: (1,3,**5**) -> (1,**1**,3)
- The player who takes the last stick loses



Minimax

```
def minimax(state):  
    v = max_value(state)  
    return action in successors(state) with value v  
  
def max_value(state):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state))  
    return v  
  
def min_value(state):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
    return v
```

Assumes opponent play **optimally**: trying to **minimize** player's value




Minimax: Analysis

- **Complete?**
 - Yes, if tree is finite
- **Time Complexity?**
 - $O(b^m)$
- **Space Complexity?**
 - $O(bm)$, with depth first exploration
- **Optimal?**
 - Yes, against optimal opponent

Chess: $b \approx 35$, $m \approx 100$ for “reasonable” games

What to do?



Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```

MAX

MIN



Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```


```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```

MAX

$\alpha = -\infty$

$\beta = \infty$

MIN



Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```


```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```

MAX

MIN

$\alpha = -\infty$
 $\beta = 3$

3



Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```


```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```

MAX

MIN

$\alpha = -\infty$
 $\beta = 3$



Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```


```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```

MAX

MIN

$\alpha = -\infty$
 $\beta = 3$




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```




Alpha-beta Pruning

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state, α, β))  
        if v >= β: return v  
        α = max(α, v)  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        if v <= α: return v  
        β = min(β, v)  
    return v
```



Alpha-beta Pruning: Analysis

- Pruning doesn't affect final result
- Good move ordering improves effectiveness of pruning
 - “Perfect ordering”: $O\left(b^{\frac{m}{2}}\right)$
- An example of **meta-reasoning**: reasoning about which computations to do first

Handling Huge/Infinite Game Trees

```
def minimax(state):  
    v = max_value(state)  
    return action in successors(state) with value v
```

```
def max_value(state):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state))  
    return v
```

```
def min_value(state):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
    return v
```

```
def minimax_with_cutoff(state):  
    v = max_value(state)  
    return action in successors(state) with value v
```

```
def max_value(state):  
    if is_cutoff(state): return eval(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state))  
    return v
```

```
def min_value(state):  
    if is_cutoff(state): return eval(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
    return v
```

is_cutoff

Return true if


- State is terminal
- Cutoff is reached (e.g., time/depth limit is exceeded)

eval

Return

- Utility for terminal states
- Heuristic value for non-terminal states

Handling Huge/Infinite Game Trees



```
def minimax_with_cutoff(state):
    v = max_value(state)
    return action in successors(state) with value v

def max_value(state):
    if is_cutoff(state): return eval(state)
    v = -∞
    for action, next_state in successors(state):
        v = max(v, min_value(next_state))
    return v

def min_value(state):
    if is_cutoff(state): return eval(state)
    v = ∞
    for action, next_state in successors(state):
        v = min(v, max_value(next_state))
    return v
```

is_cutoff

Return true if

- State is terminal
- Cutoff is reached (e.g., time/depth limit is exceeded)

eval

Return

- Utility for terminal states
- Heuristic value for non-terminal states

Evaluation Functions

How to design a good evaluation function?

- Need to understand the **problem domain**

For chess, typically linear weighted sum of features:


$$Eval(s) = w_1f_1(s) + w_2f_2(s) + \cdots + w_nf_n(s), w_i \in \mathbb{R}$$

Example features:


- Number of white queens – number of black queens
- Number of player's pieces – number of opponent's pieces

Optimizing the Search


- **Transposition table**
- **Precomputation of best moves** in the **opening** and **closing** games



Credit: thechessworld.com



Credit: chess.com



AI in Games



Credit: AAAI



Credit: IEEE Spectrum



Credit: Guardian

- **Checkers:** [Chinook](#) ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- **Chess:** [Deep Blue](#) defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- **Othello:** human champions refuse to compete against computers, who are too good.
- **Go:** [AlphaGo](#) defeated 18-time world champion Lee Sedol in 2016. Instead of using probability algorithms hard-coded by human programmers, AlphaGo uses neural networks to estimate its probability of winning.

Summary: Adversarial Search

- Games vs Search Problems:
 - Search problems: **predictable** world, deterministic and stochastic
 - Games: “**unpredictable**” opponent, strategic
- Minimax:
 - Assume **opponent behaves optimally**
 - Pick move that **maximizes** player’s **utility**
- Alpha-beta pruning
 - **Prune branches** that are useless (**will not change the decision**)
- Handling large/infinite game trees:
 - **Cutoff** the search in the middle of the game and use **evaluation function**
- Optimizing the search:
 - Transposition table: **reuse** values
 - **Precomputation** of good moves

Summary

- Informed search algorithms
 - Greedy best-first search: $f(n) = g(n)$
 - A* search: $f(n) = g(n) + h(n)$
 - Heuristics: admissibility, consistency, dominance
 - Variants of A*: IDA*, SMA*
- Local search
 - Hill climbing: pick best neighbor, repeat
 - Simulated annealing: allow bad moves from time to time, escape local optima
 - Beam search: parallel hill-climbing
 - Genetic algorithms: inspired by evolution
- Adversarial search
 - Games: opponent “unpredictable”
 - Minimax: max then min value
 - Alpha-beta pruning: prune unuseful branch
 - Handling large/infinite game trees: cutoff, eval
 - Optimizing the search: transposition table, precomputation

Coming Up Next Week

- **Machine Learning**
- **Decision Trees**
 - Entropy and Information Gain
 - Different types of attributes
 - Pruning
- **Performance Measures**
 - Precision, Recall, F1
 - ROC

To Do

- **Lecture Training 3**
 - +100 Free EXP
 - +50 Early bird bonus
- **Problem Set 1**
 - Due Saturday, 2nd September (Tomorrow)